

jxh3ztegx

January 26, 2025

```
[62]: import pandas as pd
import numpy as np
from sklearn.cluster import KMeans
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from scipy.spatial.distance import cdist
from sklearn.metrics import accuracy_score, classification_report
```

```
[26]: customers = pd.read_csv(r"C:\Users\vishn\OneDrive\Desktop\FSD\intern data_\u2192science\Customers - Customers.csv")
products = pd.read_csv(r"C:\Users\vishn\OneDrive\Desktop\FSD\intern data_\u2192science\Products - Products.csv")
transactions= pd.read_csv(r"C:\Users\vishn\OneDrive\Desktop\FSD\intern data_\u2192science\Transactions - Transactions.csv")
```

```
[12]: customers.head()
```

```
[12]:
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

```
[13]: products.head()
```

```
[13]:
```

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

```
[27]: transactions.head()
```

```
[27]: TransactionID CustomerID ProductID TransactionDate Quantity \
0      T00001      C0199      P067 2024-08-25 12:38:23      1
1      T00112      C0146      P067 2024-05-27 22:23:54      1
2      T00166      C0127      P067 2024-04-25 7:38:55      1
3      T00272      C0087      P067 2024-03-26 22:55:37      2
4      T00363      C0070      P067 2024-03-21 15:10:10      3
```

```
TotalValue Price
0      300.68 300.68
1      300.68 300.68
2      300.68 300.68
3      601.36 300.68
4      902.04 300.68
```

```
[28]: transactions['TotalSpend'] = transactions['Quantity'] * transactions['Price']
merged_data = transactions.merge(customers, on='CustomerID').merge(products,
↪on='ProductID')
```

```
[29]: customer_features = merged_data.groupby('CustomerID').agg({
      'TotalSpend': 'sum',
      'Quantity': 'sum',
      'TransactionID': 'count',
      'Region': lambda x: x.mode()[0],
      'Category': lambda x: x.mode()[0]
}).reset_index().rename(columns={'TransactionID': 'TransactionCount'})
```

```
[30]: encoder = OneHotEncoder()
encoded_categorical = encoder.fit_transform(customer_features[['Region',
↪'Category']]).toarray()
encoded_columns = encoder.get_feature_names_out(['Region', 'Category'])
```

```
[40]: encoded_df = pd.DataFrame(encoded_categorical, columns=encoded_columns)
feature_matrix = pd.concat([customer_features[['CustomerID', 'TotalSpend',
↪'Quantity', 'TransactionCount']], encoded_df], axis=1)
```

```
[41]: scaler = StandardScaler()
numeric_columns = ['TotalSpend', 'Quantity', 'TransactionCount']
feature_matrix[numeric_columns] = scaler.
↪fit_transform(feature_matrix[numeric_columns])
```

```
[42]: kmeans = KMeans(n_clusters=5, random_state=42)
feature_matrix['Cluster'] = kmeans.fit_predict(feature_matrix.
↪drop(columns=['CustomerID']))
```

C:\Users\vishn\anaconda3\new files\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`

explicitly to suppress the warning

```
super().__check_params_vs_input(X, default_n_init=10)
C:\Users\vishn\anaconda3\new files\Lib\site-
packages\sklearn\cluster\_kmeans.py:1436: UserWarning: KMeans is known to have a
memory leak on Windows with MKL, when there are less chunks than available
threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
warnings.warn(
```

```
[43]: similarity_results = []
for cluster in feature_matrix['Cluster'].unique():
    cluster_data = feature_matrix[feature_matrix['Cluster'] == cluster]
    cluster_customers = cluster_data['CustomerID']
    cluster_features = cluster_data.drop(columns=['CustomerID', 'Cluster'])
    similarity_matrix = cosine_similarity(cluster_features)
    for idx, customer_id in enumerate(cluster_customers):
        similar_customers = [
            (cluster_customers.iloc[i], similarity_matrix[idx, i])
            for i in np.argsort(similarity_matrix[idx])[-4:-1] # Top 3
        ]
        #excluding self
    similarity_results.append({
        'CustomerID': customer_id,
        'Lookalikes': [x[0] for x in similar_customers],
        'Scores': [x[1] for x in similar_customers]
    })

lookalike_df = pd.DataFrame(similarity_results)
```

```
[44]: lookalike_df.to_csv('Lookalike_KMeans.csv', index=False)
```

```
[45]: top_customers = feature_matrix.sort_values(by='TotalSpend', ascending=False).
    #head(50)['CustomerID']
merged_data['IsTopCustomer'] = merged_data['CustomerID'].apply(lambda x: 1 if x
    #in top_customers.values else 0)
```

```
[48]: merged_data = merged_data.merge(customer_features[['CustomerID',
    # 'TransactionCount']], on='CustomerID', how='left')
X = merged_data[['TotalSpend', 'Quantity', 'TransactionCount']]
y = merged_data['IsTopCustomer']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    #random_state=42)
```

```
[49]: clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)
```

```
[49]: RandomForestClassifier(random_state=42)
```

```
[50]: y_pred = clf.predict(X_test)
```

```
[51]: print("Accuracy Score:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Accuracy Score: 0.76

Classification Report:

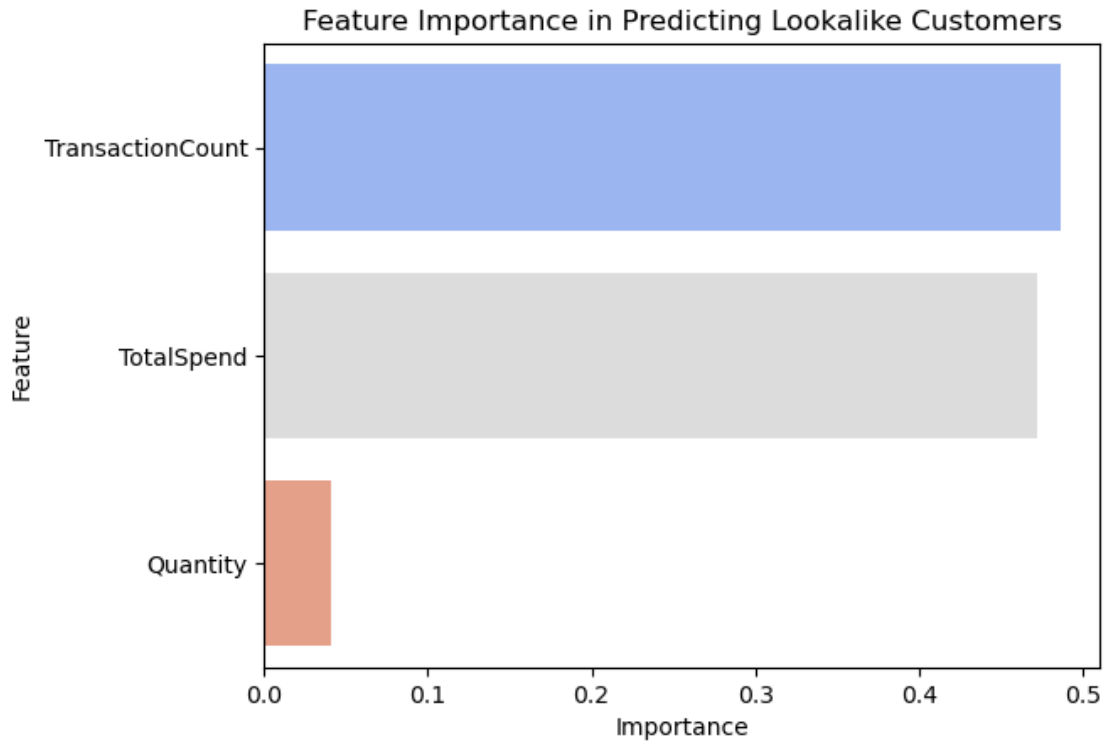
	precision	recall	f1-score	support
0	0.80	0.83	0.82	129
1	0.67	0.63	0.65	71
accuracy			0.76	200
macro avg	0.74	0.73	0.73	200
weighted avg	0.76	0.76	0.76	200

```
[54]: importances = clf.feature_importances_
feature_names = X.columns
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
    ↳ importances}).sort_values(by='Importance', ascending=False)
print("\nFeature Importance:")
print(importance_df)
```


Feature Importance:

	Feature	Importance
2	TransactionCount	0.486229
0	TotalSpend	0.472372
1	Quantity	0.041399

```
[55]: sns.barplot(x=importance_df['Importance'], y=importance_df['Feature'],
    ↳ palette='coolwarm')
plt.title('Feature Importance in Predicting Lookalike Customers')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```



```
[57]: similarity_features = feature_matrix.drop(columns=['CustomerID', 'Cluster'])
      cosine_sim = cosine_similarity(similarity_features)
      sns.heatmap(cosine_sim[:10, :10], annot=True, cmap='coolwarm',
                  xticklabels=False, yticklabels=False)
      plt.title("Similarity Matrix (Top 10 Customers)")
      plt.show()
```



```
[58]: feature_matrix.set_index('CustomerID', inplace=True)
```

```
[61]: pca = PCA(n_components=10)
reduced_features = pca.fit_transform(feature_matrix)
```

```
[63]: cosine_sim = cosine_similarity(reduced_features)
euclidean_distances = cdist(reduced_features, reduced_features,
    ↪ metric='euclidean')
```

```
[64]: combined_similarity = cosine_sim - 0.5 * (euclidean_distances / np.
    ↪ max(euclidean_distances))
```

```
[65]: lookalikes = {}
for idx, row in enumerate(combined_similarity):
    similar_customers = np.argsort(row)[-4:-1][::-1] # Top 3 most similar
    ↪ customers (excluding self)
    customer_id = feature_matrix.index[idx]
    similar_ids = [(feature_matrix.index[i], row[i]) for i in similar_customers]
    lookalikes[customer_id] = similar_ids
```

```
[66]: lookalike_df = pd.DataFrame([
    {
```

```

        'CustomerID': cust_id,
        'Lookalike1': similar_list[0][0],
        'Score1': similar_list[0][1],
        'Lookalike2': similar_list[1][0],
        'Score2': similar_list[1][1],
        'Lookalike3': similar_list[2][0],
        'Score3': similar_list[2][1],
    }
    for cust_id, similar_list in lookalikes.items()
])

print("\nTop 3 Lookalikes with Scores:")
print(lookalike_df.head())

```

Top 3 Lookalikes with Scores:

	CustomerID	Lookalike1	Score1	Lookalike2	Score2	Lookalike3	Score3
0	C0001	C0048	0.976839	C0190	0.962897	C0181	0.917362
1	C0002	C0088	0.932069	C0092	0.893227	C0040	0.655431
2	C0003	C0052	0.858355	C0031	0.837140	C0076	0.834833
3	C0004	C0087	0.910789	C0165	0.895171	C0082	0.850409
4	C0005	C0186	0.989233	C0007	0.968710	C0146	0.937738

```
[67]: lookalike_df.to_csv('Extended_Lookalike.csv', index=False)
```

```
[68]: lookalike_df.head()
```

```
[68]:
```

	CustomerID	Lookalike1	Score1	Lookalike2	Score2	Lookalike3	Score3
0	C0001	C0048	0.976839	C0190	0.962897	C0181	0.917362
1	C0002	C0088	0.932069	C0092	0.893227	C0040	0.655431
2	C0003	C0052	0.858355	C0031	0.837140	C0076	0.834833
3	C0004	C0087	0.910789	C0165	0.895171	C0082	0.850409
4	C0005	C0186	0.989233	C0007	0.968710	C0146	0.937738

```
[ ]:
```